

## Problem Set 2:

### CS525 – Advance Database Organisation

**Exercise 4.1.1:** Suppose blocks hold either three records, or ten key-pointer pairs. As a function of  $n$ , the number of records, how many blocks do we need to hold a data file and:

- a) A dense index?
- b) A sparse index?

**Answer 4.1.1:**

- a) For the dense Index we need  $n/3$  blocks are needed for records and  $n/10$  for the index and in total  $n/3 + n/10 = 13n/30$
- b) In the Sparse Index we still need  $n/3$  for the records but typically in sparse index we need only room for  $n/3$  in the index file as well. So in total  $n/30$  blocks for total of  **$11n/30$**  blocks.

**Exercise 4.3.3 :** Suppose pointers are 4 bytes long, and keys are 12 bytes long. How many keys and pointers will a block of 16,384 bytes have?

**Answer 4.3.3:**

$$4n + 12(n+1) \geq 16384$$

$$16n + 12 \geq 16384$$

$$16n \geq 16372$$

$$n \geq 1023.25$$

$$n = 1024$$

**Exercise 4.3.4:** What are the minimum numbers of keys and pointers in Btree (i) interior nodes and (M) leaves, when:

- a)  $n = 10$ ; i.e., a block holds 10 keys and 11 pointers.
- b)  $n = 11$ ; i.e., a block holds 11 keys and 12 pointers.

**Answer 4.3.4:**


- a) For the interior nodes there would be 5 keys and 6 pointers and at the Leaf Node there would be 5 keys and 5 pointers not counting the extra pointer for the next leaf node.
- b) For the interior nodes there would be 6 keys and 7 pointers and at the leaf node there would be 6 keys and 6 pointers.

**Exercise 4.3.5:** Execute the following operations on Fig. 4.23. Describe the changes for operations that modify the tree.


- a) Lookup the record with key 41.
- b) Lookup the record with key 40.
- c) Lookup all records in the range 20 to 30.
- d) Lookup all records with keys less than 30.
- e) Lookup all records with keys greater than 30.
- f) Insert a record with key 1.
- g) Insert records with keys 14 through 16.
- h) Delete the record with key 23.
- i) Delete all the records with keys 23 and higher.

**Answer 4.3.5:**


- a) Lookup 41:
  - i. As  $41 > 13$ , we will go for the right sub-tree.
  - ii. In the right subtree as  $41 > 23$  and  $41 > 31$  but  $41 < 43$ , so 41 fits exactly between  $31 < 41 < 43$  (keys).
  - iii. Then we will go for keys 31, 37 and **41**.
  - iv. We find 41 as the last key in level 3, so we return 41 key with the record with it.
- b) Lookup 40:
  - i. As  $40 > 13$ , go for the right sub-tree.
  - ii. In the right subtree at level 2 we find keys 23, 31 and 43.
  - iii. As 40 fits perfectly between  $23 < 40 < 43$ , we go for the keys 31, 37 and 41 but conclude that record with key 40 is not present at the leaf node.
- c) Lookup all records in the range 20 to 30:
  - i. As 20 and 30 both are  $> 13$  so we go for the right subtree straight away.
  - ii. In the right subtree we go for keys which are less than 23 and less than 31.
  - iii. So as 13, 17 and 19 don't fall into the criteria we go for the next leaf node that is 23 and 29.
  - iv. As 23 and 29 falls in the range 20 to 30. Return the record attached to keys 23 and 29 in this case.
- d) Lookup all records with keys less than 30.:
  - i. As now the range is from 0 to 30 and 0 is less than 13, and 30 is greater than 13, so go to leaf node and see the records with keys 0 to 30 and return the data files attach with it.
  - ii. So, the data returned will be with keys: 2,3,5,7,11,13,17,19,23,29.
  - iii. Total of 10 records.
- e) Lookup all records with keys greater than 30:
  - i. Similar as the above case.
  - ii. 5 data records will be returned, they are: 31,37,41,43 and 47 as they are all greater than the key 30.
- f) Insert a record with key 1:




g) Insert records with keys 14 through 16:



h). Delete the record with key 23:



i). Delete all the records with keys 23 and higher:

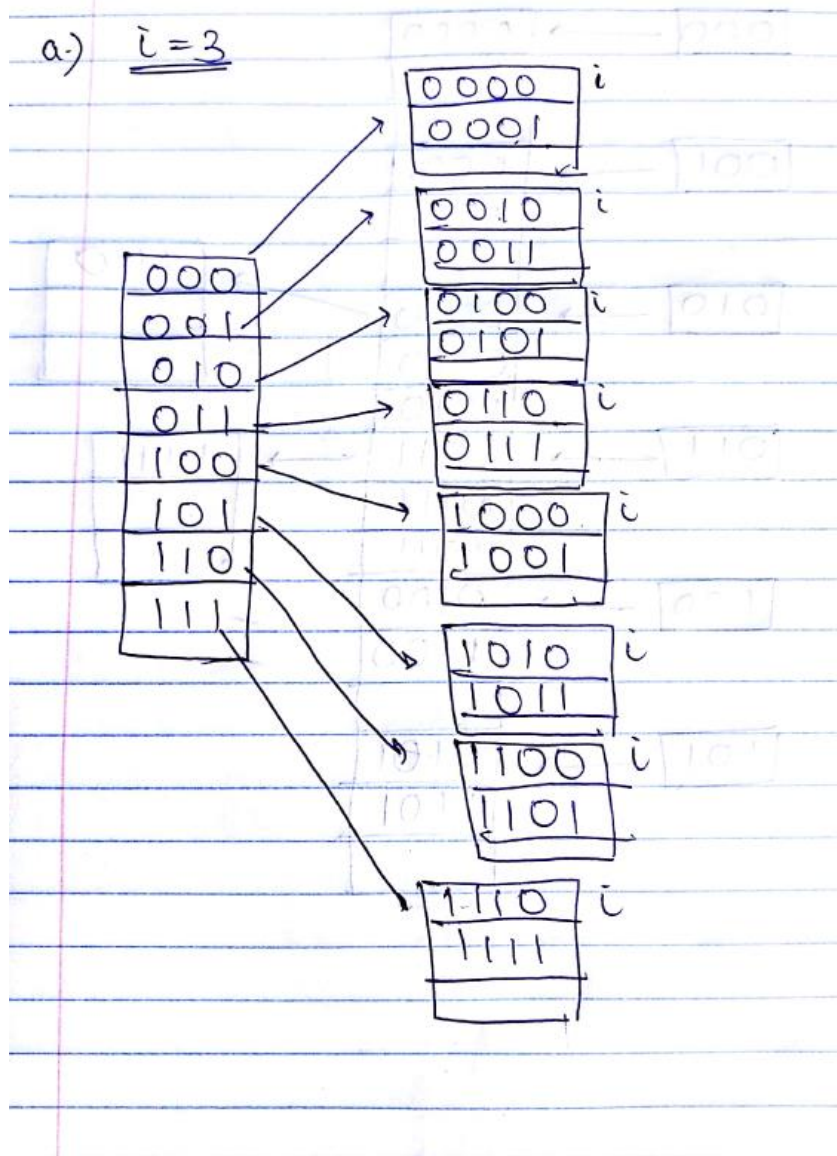


**Exercise 4.4.6:** Suppose keys are hashed to four-bit sequences, as in our examples of extensible and linear hashing in this section. However, also suppose that blocks can hold three records, rather than the two-record blocks of our examples. If we start with a hash table with two empty blocks (corresponding to 0 and 1), show the organization after we insert records with keys:

- a) 0000,0001,..., 1111, and the method of hashing is extensible hashing.
- b) 0000, 0001,..., 1111, and the method of hashing is linear hashing with a capacity threshold of 100%.
- c) 1111,1110,..., 0000, and the method of hashing is extensible hashing.
- d) 1111,1110,..., 0000, and the method of hashing is linear hashing with a capacity threshold of 75%.

**Answer 4.4.6:**

**a).**

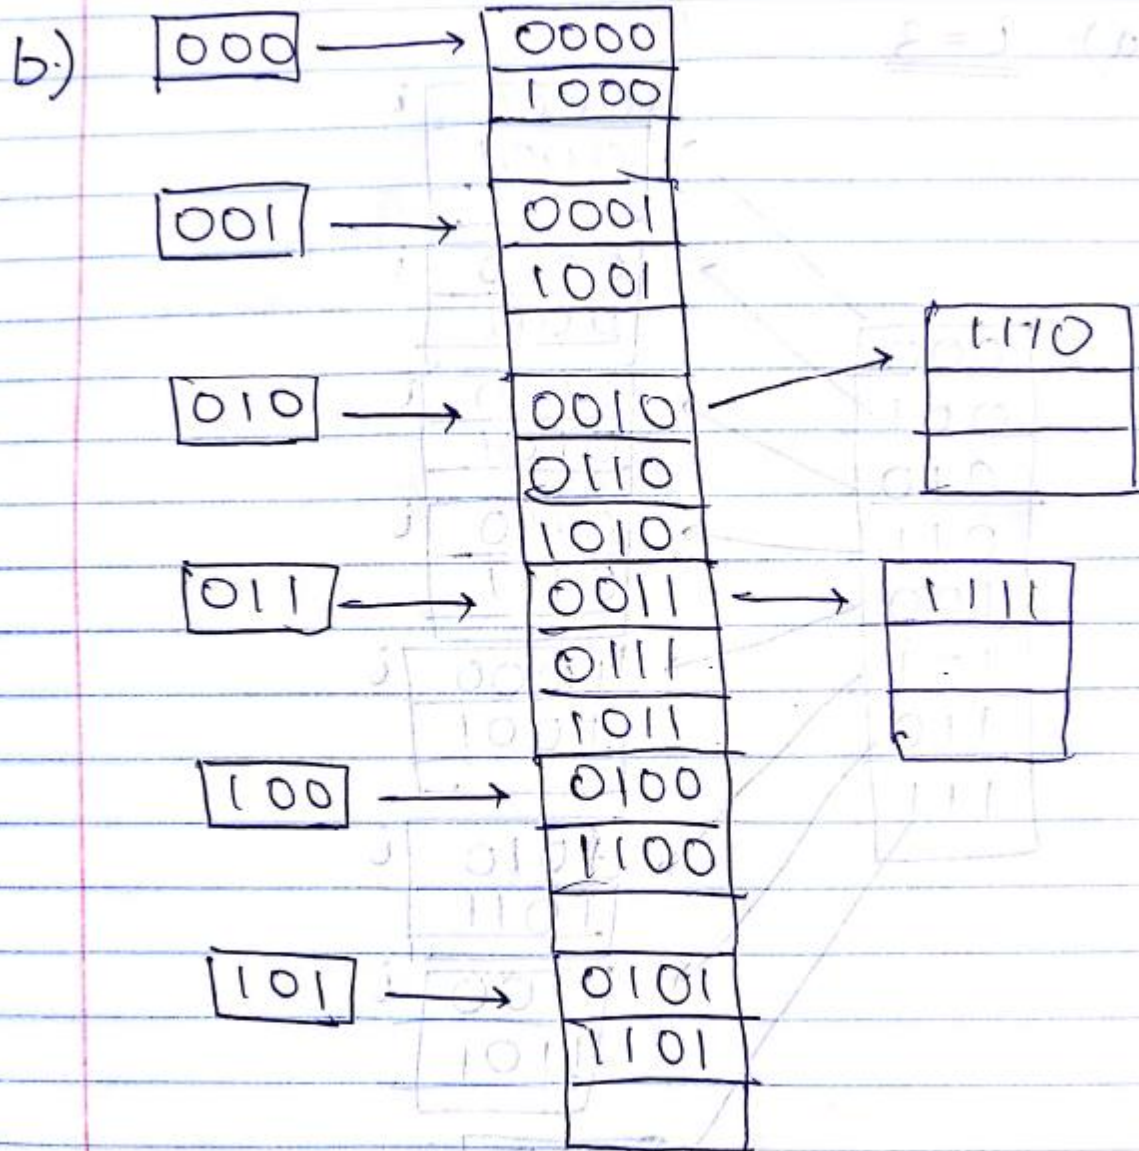




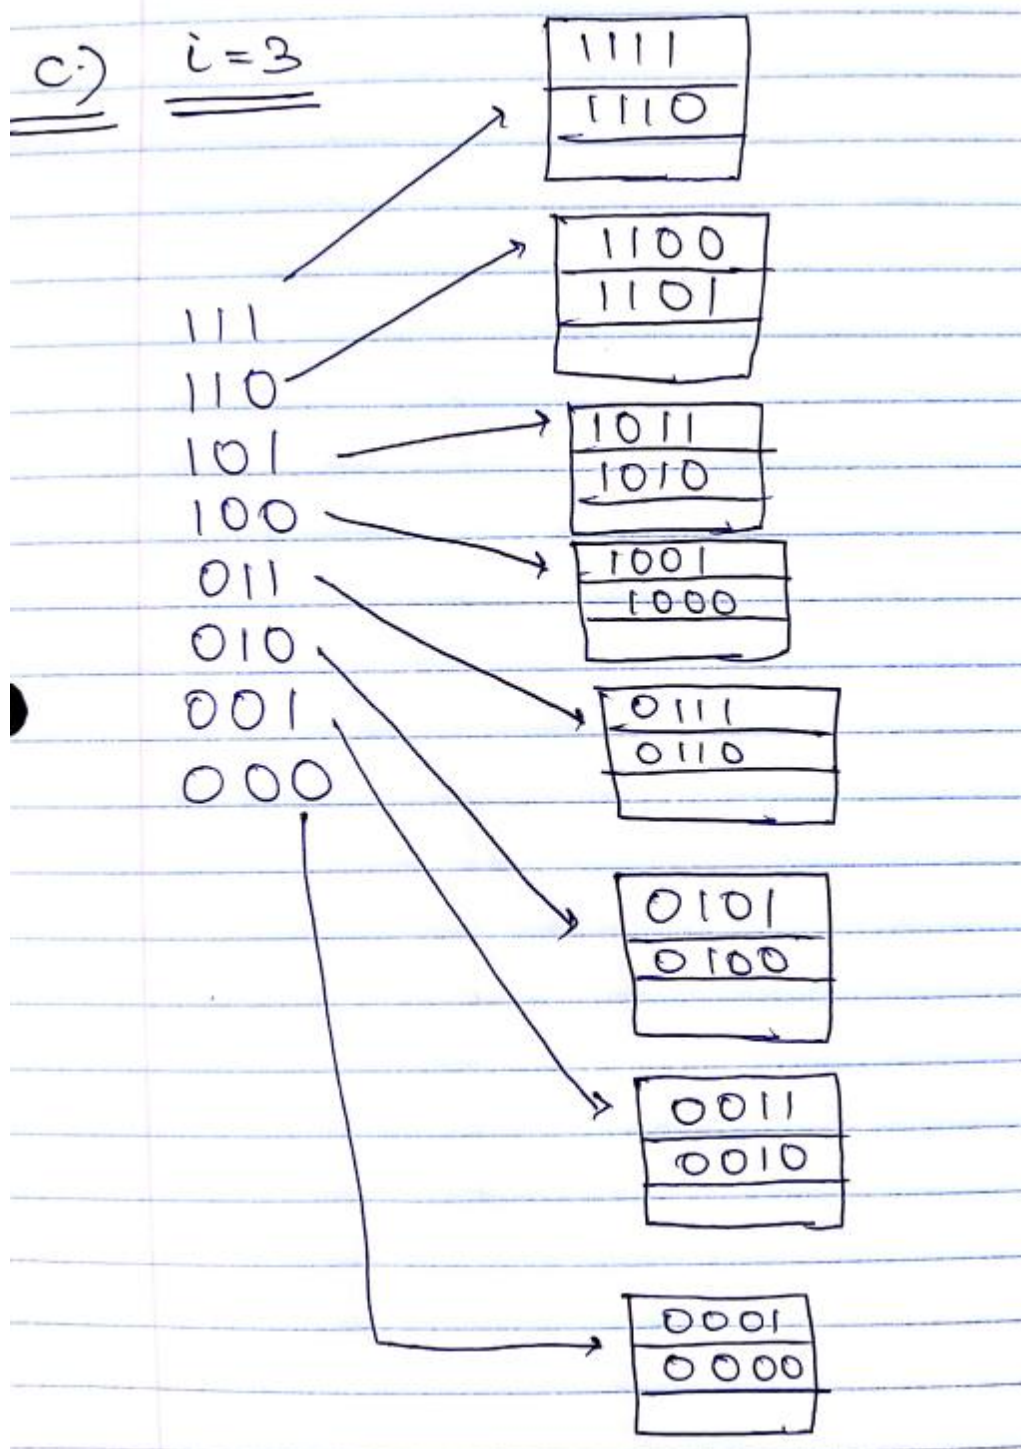
b):

$n \rightarrow 6$   
 $i \rightarrow 3$   
 $\pi \rightarrow 16$

88.89 %



c).



d).

$n \rightarrow 8$   
 $i \rightarrow 3$   
 $m \rightarrow 16$

d.)

000

0000
1000

~~000~~

001

0001
1001

010

1010
0010

011

0011
1011

100

0100
1100

101

0101
1101

110

1110
0110

111

0111
1111

66.67%